dog_app

February 9, 2021

1 Convolutional Neural Networks

1.1 Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with **'(IMPLEMENTATION)'** in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

> **Note**: Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a **'Question X'** header. Carefully read each question and provide thorough answers in the following text boxes that begin with **'Answer:'**. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

> **Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.
Step 0: Import Datasets
Make sure that you've downloaded the required human and dog datasets:
Note: if you are using the Udacity workspace, you *DO NOT* need to re-download these - they can be found in the `/data` folder as noted in the cell below.

- Download the dog dataset. Unzip the folder and place it in this project's home directory, at the location /dog_images.

- Download the human dataset. Unzip the folder and place it in the home directory, at location /lfw.

Note: If you are using a Windows machine, you are encouraged to use 7zip to extract the folder.

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human_files and dog_files.

```
In [68]: import numpy as np
         from glob import glob

         # load filenames for human and dog images
         human_files = np.array(glob("/data/lfw/*/*"))
         dog_files = np.array(glob("/data/dog_images/*/*/*"))

         # print number of images in each dataset
         print('There are %d total human images.' % len(human_files))
         print('There are %d total dog images.' % len(dog_files))
```

```
There are 13233 total human images.
There are 8351 total dog images.
```

Step 1: Detect Humans

In this section, we use OpenCV's implementation of Haar feature-based cascade classifiers to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on github. We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [69]: import cv2
         import matplotlib.pyplot as plt
         %matplotlib inline

         # extract pre-trained face detector
         face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

         # load color (BGR) image
         img = cv2.imread(human_files[0])
         # convert BGR image to grayscale
         gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

         # find faces in image
         faces = face_cascade.detectMultiScale(gray)

         # print number of faces detected in the image
         print('Number of faces detected:', len(faces))
```

```
        # get bounding box for each detected face
        for (x,y,w,h) in faces:
            # add bounding box to color image
            cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

        # convert BGR image to RGB for plotting
        cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

        # display the image, along with bounding box
        plt.imshow(cv_rgb)
        plt.show()

Number of faces detected: 1
```


Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

3

1.1.1   Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [70]: # returns "True" if face is detected in image stored at img_path
         def face_detector(img_path):
             img = cv2.imread(img_path)
             gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
             faces = face_cascade.detectMultiScale(gray)
             return len(faces) > 0
```

1.1.2   (IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.
- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?
    Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.
    Answer: (You can print out your results and/or write your percentages in this cell)

```
In [71]: from tqdm import tqdm

         human_files_short = human_files[:100]
         dog_files_short = dog_files[:100]

         #-#-# Do NOT modify the code above this line. #-#-#

         ## TODO: Test the performance of the face_detector algorithm
         ## on the images in human_files_short and dog_files_short.
         human_file_faces_detected = 0
         for human in human_files_short:
             detected = face_detector(human)
             if detected:
                 human_file_faces_detected += 1

         dog_file_faces_detected = 0
         for dog in dog_files_short:
             detected = face_detector(dog)
             if detected:
                 dog_file_faces_detected += 1

         print(f"Percentage of human face detection in human_files at {human_file_faces_detected
         print(f"Percentage of human face detection in dog_files at {dog_file_faces_detected}%")
```

```
Percentage of human face detection in human_files at 98%
Percentage of human face detection in dog_files at 17%
```

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [72]: ### (Optional)
         ### TODO: Test performance of anotherface detection algorithm.
         ### Feel free to use as many code cells as needed.
```

Step 2: Detect Dogs
In this section, we use a pre-trained model to detect dogs in images.

1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on ImageNet, a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of 1000 categories.

```
In [73]: import torch
         import torchvision.models as models

         # define VGG16 model
         VGG16 = models.vgg16(pretrained=True)

         # check if CUDA is available
         use_cuda = torch.cuda.is_available()

         # move model to GPU if CUDA is available
         if use_cuda:
             VGG16 = VGG16.cuda()
```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as `'dogImages/train/001.Affenpinscher/Affenpinscher_00001.jpg'`) as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the PyTorch documentation.

```python
In [74]: from PIL import Image
         import torchvision.transforms as transforms

         def VGG16_predict(img_path):
             '''
             Use pre-trained VGG-16 model to obtain index corresponding to
             predicted ImageNet class for image at specified path

             Args:
                 img_path: path to an image

             Returns:
                 Index corresponding to VGG-16 model's prediction
             '''

             ## TODO: Complete the function.
             ## Load and pre-process an image from the given img_path
             ## Return the *index* of the predicted class for that image

             img = Image.open(img_path)

             transform = transforms.Compose([
                 transforms.Resize(256),
                 transforms.CenterCrop(224),
                 transforms.ToTensor(),
                 transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225))
             ])

             image = transform(img).unsqueeze(0)

             if use_cuda:
                 image = image.cuda()

             output = VGG16(image)

             if use_cuda:
                 output = output.cpu()

             index = output.data.numpy().argmax()
             return index # predicted class index

In [75]: print(VGG16_predict('/data/dog_images/train/001.Affenpinscher/Affenpinscher_00001.jpg')

252
```

1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the dictionary, you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog_detector function below, which returns True if a dog is detected in an image (and False if not).

```python
In [76]: ### returns "True" if a dog is detected in the image stored at img_path
         def dog_detector(img_path):
             ## TODO: Complete the function.
             index = VGG16_predict(img_path)
             if 150 < index < 269:
                 return True
             else:
                 return False
```

1.1.6 (IMPLEMENTATION) Assess the Dog Detector

Question 2: Use the code cell below to test the performance of your dog_detector function.
- What percentage of the images in human_files_short have a detected dog?
- What percentage of the images in dog_files_short have a detected dog?
  Answer:

```python
In [77]: ### TODO: Test the performance of the dog_detector function
         ### on the images in human_files_short and dog_files_short.
         human_file_dog_detected = 0
         for human in human_files_short:
             detected = dog_detector(human)
             if detected:
                 human_file_dog_detected += 1

         dog_file_dog_detected = 0
         for dog in dog_files_short:
             detected = dog_detector(dog)
             if detected:
                 dog_file_dog_detected += 1

         print(f"Percentage of dog detection in human_files at {human_file_dog_detected}%")
         print(f"Percentage of dog detection in dog_files at {dog_file_dog_detected}%")

Percentage of dog detection in human_files at 2%
Percentage of dog detection in dog_files at 99%
```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as Inception-v3, ResNet-50, etc). Please use

the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [78]:  ### (Optional)
          ### TODO: Report the performance of another pre-trained network.
          ### Feel free to use as many code cells as needed.
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet*!), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany   Welsh Springer Spaniel

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever   American Water Spaniel

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador   Chocolate Labrador

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imabalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

1.1.7   (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate data loaders for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find this documentation on custom datasets to be a useful resource. If you

are interested in augmenting your training and/or validation data, check out the wide variety of transforms!

```
In [91]: ### TODO: Write data loaders for training, validation, and test sets
         ## Specify appropriate transforms, and batch_sizes
         # prepare data loaders
         from torchvision import datasets

         # how many samples per batch to load
         batch_size = 20
         num_workers = 0

         #transforms
         training_transformers = transforms.Compose([transforms.Resize(256),
                                                  transforms.RandomHorizontalFlip(),
                                                  transforms.CenterCrop(224),
                                                  transforms.ToTensor(),
                                                  transforms.Normalize([0.5, 0.5, 0.5], [0.5,

         transformers = transforms.Compose([transforms.Resize(256),
                                         transforms.CenterCrop(224),
                                         transforms.ToTensor(),
                                         transforms.Normalize([0.5, 0.5, 0.5], [0.5, 0.5, 0.5

         # load the dataset
         train_data = datasets.ImageFolder(root='/data/dog_images/train', transform=training_tra
         validation_data = datasets.ImageFolder(root='/data/dog_images/valid', transform=transfo
         test_data = datasets.ImageFolder(root='/data/dog_images/test', transform=transformers)

         train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size, num_worke
         validation_loader = torch.utils.data.DataLoader(validation_data, batch_size=batch_size,
         test_loader = torch.utils.data.DataLoader(test_data, batch_size=batch_size, num_workers

         loaders_scratch = {"train": train_loader, "valid": validation_loader, "test": test_load

In [92]: print('Num training images: ', len(train_data))
         print('Num test images: ', len(test_data))
         #print(train_data.classes)
         output_nodes = len(train_data.classes)
         print(output_nodes)

Num training images:  6680
Num test images:  836
133
```

Question 3: Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and

why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

    Answer: I resized and cropped all photos to ensure they were the same size. I also took the center of the photos to get a clearer focus of the image.

1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```python
In [93]: import torch.nn as nn

         # define the CNN architecture
         class Net(nn.Module):
             ### TODO: choose an architecture, and complete the class
             def __init__(self):
                 super(Net, self).__init__()
                 self.conv1 = nn.Conv2d(3, 16, 3, padding=1)
                 # convolutional layer (sees 16x16x16 tensor)
                 self.conv2 = nn.Conv2d(16, 32, 3, padding=1)
                 # max pooling layer
                 self.pool = nn.MaxPool2d(2, 2)
                 # Linear layer
                 # the 28 here is derived from printing x.size down below
                 self.fc1 = nn.Linear(32 * 56 * 56, 500)
                 # linear layer (500 -> output nodes) - number of different classes
                 self.fc2 = nn.Linear(500, output_nodes)
                 # dropout layer (p=0.5)
                 self.dropout = nn.Dropout(0.5)

             def forward(self, x):
                 # add sequence of convolutional and max pooling layers
                 x = self.pool(F.relu(self.conv1(x)))
                 x = self.pool(F.relu(self.conv2(x)))
                 # flatten image input
                 #print(x.size())
                 x = x.view(x.size(0), -1)
                 # add dropout layer
                 x = self.dropout(x)
                 # add 1st hidden layer, with relu activation function
                 x = F.relu(self.fc1(x))
                 # add dropout layer~
                 x = self.dropout(x)
                 # add 2nd hidden layer, with relu activation function
                 x = self.fc2(x)

                 return x

         #-#-# You do NOT have to modify the code below this line. #-#-#
```

10

```
# instantiate the CNN
model_scratch = Net()

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()
```

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

Answer: I started by adding 3 convolution layers with maxpooling and a drop out for overfitting. And then i create a forward pass by flattening a image and then adding hidden layers.

1.1.9   (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a loss function and optimizer. Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```
In [94]: import torch.optim as optim

         ### TODO: select loss function
         criterion_scratch = nn.CrossEntropyLoss()

         ### TODO: select optimizer
         optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.01)
```

1.1.10   (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. Save the final model parameters at filepath `'model_scratch.pt'`.

```
In [95]: from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
             """returns trained model"""
             # initialize tracker for minimum validation loss
             valid_loss_min = np.Inf

             for epoch in range(1, n_epochs+1):
                 # initialize variables to monitor training and validation loss
                 train_loss = 0.0
                 valid_loss = 0.0

                 ###################
                 # train the model #
                 ###################
                 model.train()
                 for batch_idx, (data, target) in enumerate(loaders['train']):
```

```python
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## find the loss and update the model parameters accordingly
        ## record the average training loss, using something like
        ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_lo

        # clear the gradients of all optimized variables
        optimizer.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model paramet
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer.step()
        # update running training loss
        train_loss += ((1 / (batch_idx + 1)) * (loss.data - train_loss))

    ######################
    # validate the model #
    ######################
    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['valid']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update running validation loss
        valid_loss += loss.item() * data.size(0)


    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
        ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss <= valid_loss_min:
        print('Validation loss decreased ({:.6f} --> {:.6f}).  Saving model ...'.fo
        torch.save(model.state_dict(), save_path)
```

```
                valid_loss_min = valid_loss

        # return trained model
        return model


    # train the model
    model_scratch = train(15, loaders_scratch, model_scratch, optimizer_scratch,
                          criterion_scratch, use_cuda, 'model_scratch.pt')
```

```
Epoch: 1        Training Loss: 4.874211         Validation Loss: 4031.181304
Validation loss decreased (inf --> 4031.181304).  Saving model ...
Epoch: 2        Training Loss: 4.760890         Validation Loss: 3878.004186
Validation loss decreased (4031.181304 --> 3878.004186).  Saving model ...
Epoch: 3        Training Loss: 4.588807         Validation Loss: 3742.001698
Validation loss decreased (3878.004186 --> 3742.001698).  Saving model ...
Epoch: 4        Training Loss: 4.435779         Validation Loss: 3665.223975
Validation loss decreased (3742.001698 --> 3665.223975).  Saving model ...
Epoch: 5        Training Loss: 4.324605         Validation Loss: 3605.954015
Validation loss decreased (3665.223975 --> 3605.954015).  Saving model ...
Epoch: 6        Training Loss: 4.233980         Validation Loss: 3566.975288
Validation loss decreased (3605.954015 --> 3566.975288).  Saving model ...
Epoch: 7        Training Loss: 4.145388         Validation Loss: 3540.696216
Validation loss decreased (3566.975288 --> 3540.696216).  Saving model ...
Epoch: 8        Training Loss: 4.055150         Validation Loss: 3487.018301
Validation loss decreased (3540.696216 --> 3487.018301).  Saving model ...
Epoch: 9        Training Loss: 3.970500         Validation Loss: 3484.558232
Validation loss decreased (3487.018301 --> 3484.558232).  Saving model ...
Epoch: 10       Training Loss: 3.895101         Validation Loss: 3432.207944
Validation loss decreased (3484.558232 --> 3432.207944).  Saving model ...
Epoch: 11       Training Loss: 3.805208         Validation Loss: 3445.767250
Epoch: 12       Training Loss: 3.722965         Validation Loss: 3408.909456
Validation loss decreased (3432.207944 --> 3408.909456).  Saving model ...
Epoch: 13       Training Loss: 3.622405         Validation Loss: 3415.900207
Epoch: 14       Training Loss: 3.500602         Validation Loss: 3410.825390
Epoch: 15       Training Loss: 3.412688         Validation Loss: 3431.195185
```

In [96]: #load the model that got the best validation accuracy
         model_scratch.load_state_dict(torch.load('model_scratch.pt'))

1.1.11   (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

In [97]: import torch.nn.functional as F

         def test(loaders, model, criterion, use_cuda):

```python
            # monitor test loss and accuracy
            test_loss = 0.
            correct = 0.
            total = 0.

            model.eval()
            for batch_idx, (data, target) in enumerate(loaders['test']):
                # move to GPU
                if use_cuda:
                    data, target = data.cuda(), target.cuda()
                # forward pass: compute predicted outputs by passing inputs to the model
                output = model(data)
                # calculate the loss
                loss = criterion(output, target)
                # update average test loss
                test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
                # convert output probabilities to predicted class
                pred = output.data.max(1, keepdim=True)[1]
                # compare predictions to true label
                correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
                total += data.size(0)

            print('Test Loss: {:.6f}\n'.format(test_loss))

            print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
                100. * correct / total, correct, total))

        # call test function
        test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

```
Test Loss: 4.040317


Test Accuracy: 10% (85/836)
```

Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

1.1.12   (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate data loaders for the training, validation, and test datasets of dog images (located at `dogImages/train`, `dogImages/valid`, and `dogImages/test`, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [98]: ## TODO: Specify data loaders
         loaders_transfer = {"train": train_loader, "valid": validation_loader, "test": test_loa
```

1.1.13   (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```
In [99]: import torchvision.models as models
         import torch.nn as nn

         ## TODO: Specify model architecture
         # define VGG16 model
         model_transfer = models.vgg16(pretrained=True)

         # Freeze training for all "features" layers
         for param in model_transfer.features.parameters():
             param.requires_grad = False

         n_inputs = model_transfer.classifier[6].in_features

         # add last linear layer
         # new layers automatically have requires_grad = True
         final_layer = nn.Linear(n_inputs, output_nodes)

         model_transfer.classifier[6] = final_layer

         if use_cuda:
             model_transfer = model_transfer.cuda()
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: I use VGG models convolution layers to do the training and freeze training for all the features layers to then have one final layer.

1.1.14   (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a loss function and optimizer. Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [100]: criterion_transfer = nn.CrossEntropyLoss()
          optimizer_transfer = optim.SGD(model_transfer.classifier.parameters(), lr=0.001)
```

1.1.15   (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. Save the final model parameters at filepath `'model_transfer.pt'`.

```
In [102]: # train the model
          n_epochs = 5
          model_transfer = train(n_epochs, loaders_transfer, model_transfer, optimizer_transfer,

          # load the model that got the best validation accuracy (uncomment the line below)
          model_transfer.load_state_dict(torch.load('model_transfer.pt'))
```

```
Epoch: 1         Training Loss: 2.654603         Validation Loss: 1353.358013
Validation loss decreased (inf --> 1353.358013).  Saving model ...
Epoch: 2         Training Loss: 1.430113         Validation Loss: 814.966197
Validation loss decreased (1353.358013 --> 814.966197).  Saving model ...
Epoch: 3         Training Loss: 1.017480         Validation Loss: 641.911969
Validation loss decreased (814.966197 --> 641.911969).  Saving model ...
Epoch: 4         Training Loss: 0.827123         Validation Loss: 561.268923
Validation loss decreased (641.911969 --> 561.268923).  Saving model ...
Epoch: 5         Training Loss: 0.723859         Validation Loss: 506.672967
Validation loss decreased (561.268923 --> 506.672967).  Saving model ...
```

1.1.16   (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [103]: test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

```
Test Loss: 0.623740
```

```
Test Accuracy: 82% (688/836)
```

1.1.17   (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [104]: ### TODO: Write a function that takes a path to an image as input
          ### and returns the dog breed that is predicted by the model.

          # list of class names by index, i.e. a name can be accessed like class_names[0]
          class_names = [item[4:].replace("_", " ") for item in train_data.classes]

          def predict_breed_transfer(img_path):
              # load the image and return the predicted breed
              img = Image.open(img_path)

              transform = transforms.Compose([transforms.Resize(256),
                                              transforms.CenterCrop(224),
```

Sample Human Output

```
                                        transforms.ToTensor(),
                                        transforms.Normalize([0.5, 0.5, 0.5], [0.5, 0.5, 0

        image = transform(img).unsqueeze(0)

        if use_cuda:
            image = image.cuda()

        output = model_transfer(image)

        if use_cuda:
            output = output.cpu()

        index = output.data.numpy().argmax()

        breed = class_names[index]

        return breed
```

Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

1.1.18   (IMPLEMENTATION) Write your Algorithm

```
In [105]: ### TODO: Write your algorithm.
          ### Feel free to use as many code cells as needed.
```

17

```
def run_app(img_path):
    ## handle cases for a human face, dog, and neither
    if face_detector(img_path):
        breed = predict_breed_transfer(img_path)
        return f"Hello Human. You look like a {breed}"
    elif dog_detector(img_path):
        breed = predict_breed_transfer(img_path)
        return f"Hello Dog.  We think this is a {breed}"
    else:
        return "No dog or human found here"
```

Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

1.1.19   (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: (Three possible points for improvement) the output is good. 1) we could have one detector to tell if its a dog or human 2) the detector could return a percentage accuracy instead of true or false 3) it could return back images with the prediction

```
In [107]:  ## TODO: Execute your algorithm from Step 6 on
           ## at least 6 images on your computer.
           ## Feel free to use as many code cells as needed.

           root = "/home/workspace/dog_project/chloe-images/"

           human_files = np.array(glob(root + "human/*"))
           dog_files = np.array(glob(root + "dog/*"))

           #print(human_files)

           ## suggested code, below
           for file in np.hstack((human_files[:3], dog_files[:3])):
               #print(file)
               img = cv2.imread(file)
               cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
               plt.imshow(cv_rgb)
               plt.show()
               print(run_app(file))
```
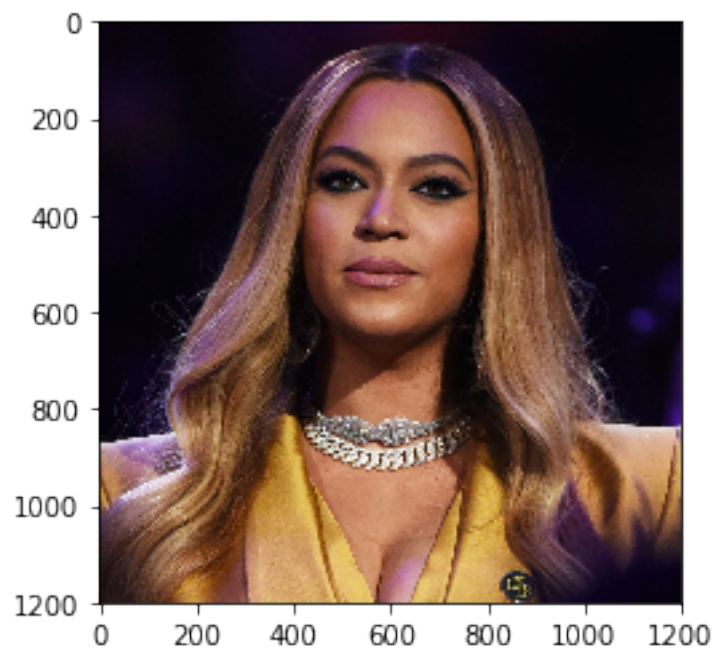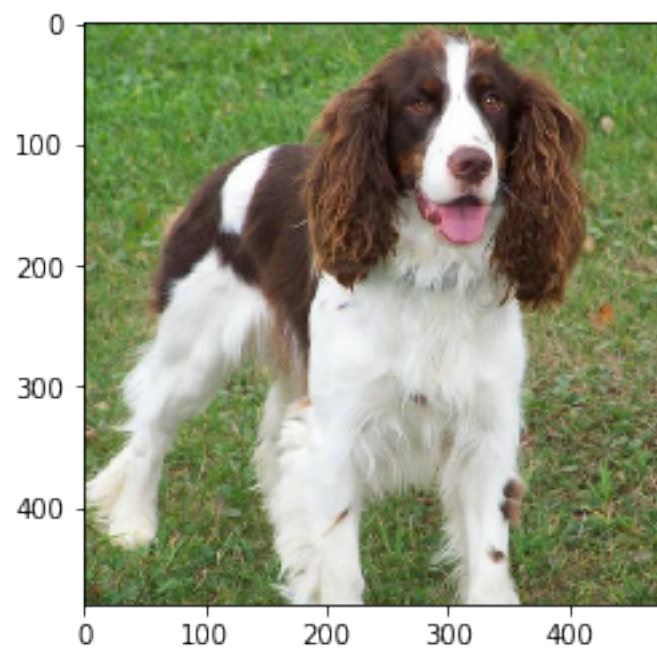
Hello Human. You look like a Afghan hound

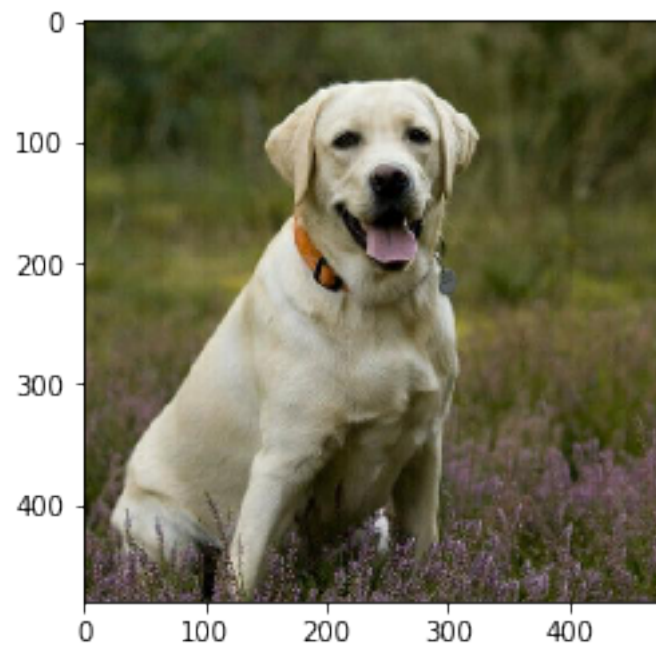Hello Human. You look like a Dachshund

Hello Human. You look like a Borzoi

Hello Dog. We think this is a English springer spaniel

Hello Dog. We think this is a Bearded collie

Hello Dog. We think this is a Labrador retriever

In [30]: import os
         os.getcwd()

Out[30]: '/home/workspace/dog_project'

In [ ]: